

# 350c

## Goodness of Fit Measures for Binary Outcome Models

Justin Grimmer (Jens Hainmueller's Slides)

Stanford

# Outline

- 1 Binary Outcomes
- 2 Calibration Plots
- 3 Discrimination
  - Classification accuracy
  - Sensitivity and specificity
- 4 Summary Accuracy Measures

# Outline

- 1 Binary Outcomes
- 2 Calibration Plots
- 3 Discrimination
  - Classification accuracy
  - Sensitivity and specificity
- 4 Summary Accuracy Measures

# Regression Models for Binary Outcomes

- Often we are interested in binary outcomes  $Y$  that take on two values 1 or 0 depending on whether an event of interest occurred or not (vote or not, war or peace, etc).
- Three widely used regression models for binary outcomes:
  - Linear Probability Model (LPM)
  - Probit regression
  - Logistic regression
- Goals
  - **Model** the probability of the event occurring as a function of the independent variables
  - **Estimate** the probability of the event occurring given certain values of the independent variables
  - **Predict** how the probability of the event occurring changes for different values of the independent variables
  - **Classify** the observations into categories based on the probability of the event occurring

# Assessing Goodness of Fit

- How do we validate that our fitted model is any good? How well does our model fit the data?
- Two core concepts to assess model fit for binary outcome models:
- **Calibration**
  - A model is well calibrated if the observed risk matches the predicted risk (e.g. for observations with a predicted probability of .2 there should be approximately 20% of  $Y = 1$  observed outcomes)
  - Well calibrated model ensures that in future samples the observed proportions of positives will be close to the estimated probabilities
- **Discrimination**
  - A model has high discrimination ability if it allows us to discriminate between low and high risk observations
  - This means observations with  $Y = 1$  should have high and observations with  $Y = 0$  should have low predicted probabilities  $\hat{Y}$
  - But we don't want a lot of  $\hat{Y}$  close to .5 if outcomes are all  $Y = 1$  or  $Y = 0$  (sharpness)

# Outline

- 1 Binary Outcomes
- 2 Calibration Plots**
- 3 Discrimination
  - Classification accuracy
  - Sensitivity and specificity
- 4 Summary Accuracy Measures

# Calibration

- In a well calibrated model, the observed risk matches the predicted risk such that in future samples the observed proportion of positives will match the predicted probabilities
- Can check this for any model prediction using calibration plots:
  - Bin predicted probabilities  $\hat{Y}$  into 10 equal interval bins  $g_k \in ([0, .1], (.1, .2], \dots, (.9, 1])$
  - In each bin  $g_k$  compute the average predicted probability  $\bar{\hat{Y}}_{\in g_k}$  and the fraction of observed positives  $\bar{Y}_{\in g_k}$
  - Plot the two vectors of averages against each other
- If the model is well calibrated, the binned averages should trace the identity line

# Calibration Plot: Correct Model

R Code

```

> n <- 1000
> x <- rnorm(n)
> xb <- 2*x
> pr <- exp(xb)/(1+exp(xb))
> y <- 1*(runif(n) < pr)
> d <- data.frame(y,x)
>
> logit <- glm(y~x, data=d, family=binomial(link="logit"))
> summary(logit)

```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-0.05078	0.08164	-0.622	0.534
x	1.99622	0.12797	15.599	<2e-16 ***


```

>
> # cut into 10 bins
> d$predprobs <- logit$fitted
> d$pptbin <- cut(d$predprobs,seq(from=0,to=1,by=0.1),include.lowest = TRUE)
> # compute average predicted probs
> rel.logit <- (d %>% group_by(pptbin) %>% summarise(out = mean(y)))
> # compute observed averages
> mpv <- (d %>% group_by(pptbin) %>% summarise(out = mean(predprobs)))
>
> pdf <- data.frame(fop = rel.logit$out,
+                   mpv = mpv$out)
>
> ggplot(pdf,aes(x=mpv,y=fop)) + geom_point(color="dodgerblue") + geom_line(color="dodgerblue") +
+   theme_economist(dkpanel = T) + geom_abline(intercept = 0,slope=1,linetype=2) +
+   xlim(0,1) + ylim(0,1) + xlab("Mean Predicted Value") + ylab("Fraction of Positives")

```



# Calibration Plot: Correct Model



# Calibration Plot: Miss-specified Model


R Code

```

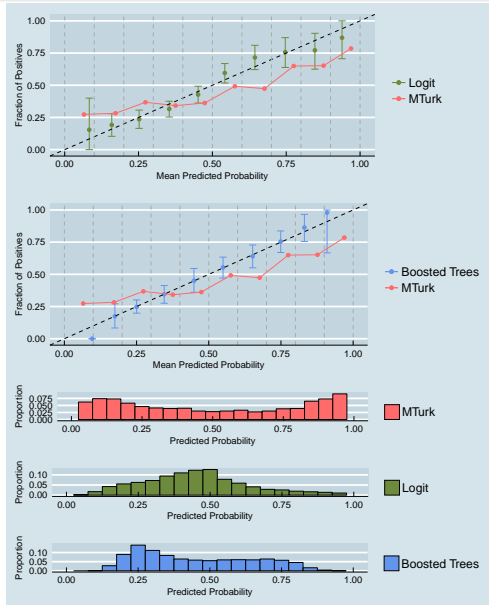
> xb <- -2*x + 1.25*x^2
> pr <- exp(xb)/(1+exp(xb))
> y <- 1*(runif(n) < pr)
> d <-data.frame(y,x)
>
> logit <- glm(y~x, data=d, family=binomial(link="logit"))
> summary(logit)
Coefficients:
              Estimate Std. Error z value Pr(>|z|)
(Intercept)  0.61150    0.07215   8.475  <2e-16 ***
x            -0.92570    0.08506  -10.883  <2e-16 ***
> d$predprobs <- logit$fitted
> head(d)
   y      x predprobs
1  1  0.07894658 0.6314440
2  0 -0.21551746 0.6923231
3  0 -0.25559249 0.7001684
4  1  0.77489539 0.4735684
5  0  2.62295669 0.1398448
6  0  0.50200285 0.5366323
>
> # cut predicted probabilities into 10 bins
> d$pptbin <- cut(d$predprobs,seq(from=0,to=1,by=0.1),include.lowest = TRUE)
> rel.logit <- (d %>% group_by(pptbin) %>% summarise(out = mean(y)))
> mpv      <- (d %>% group_by(pptbin) %>% summarise(out = mean(predprobs)))
>
> pdf <- data.frame(fop = rel.logit$out,
+                   mpv = mpv$out)
> ggplot(pdf,aes(x=mpv,y=fop)) + geom_point(color="dodgerblue") + geom_line(color="dodgerblue") +
+   theme_economist(dkpanel = T) + geom_abline(intercept = 0,slope=1,linetype=2) +
+   xlim(0,1) + ylim(0,1) + xlab("Mean Predicted Value") + ylab("Fraction of Positives")

```

# Example with Miss-specified Model



# Predicting Recidivism



# Outline

- 1 Binary Outcomes
- 2 Calibration Plots
- 3 Discrimination
  - Classification accuracy
  - Sensitivity and specificity
- 4 Summary Accuracy Measures

# Discrimination

- Good calibration of a model is not sufficient for model validation
  - A silly model that predicts the same predicted probability for every observation will have good calibration (in future samples the observed proportion will be close to the predicted probability)
  - However, such a model would not be helpful since it does not help us to discriminate between high and low risk observations
- Need to check **discrimination ability** of a model
  - Discrimination is high if  $Y = 1$  observations correspond with high  $\hat{Y}$  values and observations with  $Y = 0$  correspond to low  $\hat{Y}$  values
- Discrimination can be captured in various ways including classification accuracy, sensitivity, specificity, precision, recall, etc.
- Many metrics are used and different fields have different preferences

# Outline

- 1 Binary Outcomes
- 2 Calibration Plots
- 3 Discrimination
  - Classification accuracy
  - Sensitivity and specificity
- 4 Summary Accuracy Measures

# Classification accuracy

For classification we choose a cutoff value on the probability scale, say  $c = .5$ , and classify all predicted values above  $c$  as predicting a positive, and all below  $c$  as predicting a negative

The **contingency table** or **confusion matrix** for a binary classifier is given by:

	Observed	
	Positive ( $Y = 1$ )	Negative ( $Y = 0$ )
Predicted positive ( $\hat{Y} > c$ )	True Positive (TP)	False Positive (FP)
Predicted negative ( $\hat{Y} < c$ )	False Negative (FN)	True Negative (TN)
	Total Positive (P)	Total Negative (N)

- **Accuracy:**  $ACC = (TP + TN)/(P + N)$ 
  - Percent correctly classified of all predictions



# Recidivism Example: Classification accuracy

R Code

```

> source("http://pcwww.liv.ac.uk/~william/R/crosstab.r")
> d <- read.csv("bail.csv")
> d$y <- d$two_year_recid
> crosstab(d,row.vars="y")

y      Count Total %
0      524.0   52.4
1      476.0   47.6
Sum 1000.0  100.0
>
> # model 1
> logit1 <- glm(y~sex+age+as.factor(race),data=d)
> summary(logit1)

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)    0.653519   0.055209   11.837 < 2e-16 ***
sex             -0.162660   0.038255   -4.252 2.32e-05 ***
age             -0.006605   0.001296   -5.096 4.16e-07 ***
as.factor(race)2  0.162315   0.033086    4.906 1.09e-06 ***
as.factor(race)3 -0.012272   0.057799   -0.212  0.832
as.factor(race)4 -0.195496   0.183733   -1.064  0.288
as.factor(race)5 -0.475177   0.480880   -0.988  0.323
---
> d$yhat1 <- logit1$fitted
> d$ypred1 <- as.numeric(d$yhat1>.5)

```

# Recidivism Example: Classification accuracy

R Code

```
> # model 2  
> logit2 <- glm(y~sex+age+as.factor(race)+log(1+priors_count),data=d)  
> summary(logit2)
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	0.553299	0.053324	10.376	< 2e-16 ***
sex	-0.104176	0.036774	-2.833	0.00471 **
age	-0.007815	0.001237	-6.318	3.99e-10 ***
as.factor(race)2	0.089658	0.032197	2.785	0.00546 **
as.factor(race)3	-0.021096	0.054915	-0.384	0.70094
as.factor(race)4	-0.070564	0.174957	-0.403	0.68680
as.factor(race)5	-0.573363	0.456928	-1.255	0.20984
log(1 + priors_count)	0.166681	0.016016	10.407	< 2e-16 ***

---

```
> d$yhat2 <- logit2$fitted  
> d$ypred2 <- as.numeric(d$yhat2>.5)
```

# Recidivism Example: Classification accuracy

R Code

```

> # model 1
> crosstab(d,row.vars="ypred1",col.vars="y",type="frequency",percentages=F)
      y      0      1  Sum
ypred1
0      352    204    556
1      172    272    444
Sum     524    476   1000
> ACCm1 <- mean(d$ypred1 == d$y)
> ACCm1
[1] 0.624
>
> # model 2
> crosstab(d,row.vars="ypred2",col.vars="y",type="frequency",percentages=F)
      y      0      1  Sum
ypred2
0      378    178    556
1      146    298    444
Sum     524    476   1000
> ACCm2 <- mean(d$ypred2 == d$y)
> ACCm2
[1] 0.676
>
> # all negative model (null model)
> ACCneg <- mean(0 == d$y)
> ACCneg
[1] 0.524
>
> # percent increase in accuracy over null model
> ((ACCm2-ACCneg)/ACCneg)*100
[1] 29.00763

```

# Classification accuracy (Cutpoint .85)

R Code

```
> # model 1
> d$ypred1 <- as.numeric(d$yhat1>.85)
>
> # model 2
> d$ypred2 <- as.numeric(d$yhat2>.85)
>
> # model 1
> crosstab(d,row.vars="ypred1",col.vars="y",type="frequency",percentages=F)
      y      0      1    Sum
ypred1
0      524    476  1000
Sum    524    476  1000
> ACCm1 <- mean(d$ypred1 == d$y)
> ACCm1
[1] 0.524
>
> # model 2
> crosstab(d,row.vars="ypred2",col.vars="y",type="frequency",percentages=F)
      y      0      1    Sum
ypred2
0      521    455   976
1         3     21    24
Sum    524    476  1000
> ACCm2 <- mean(d$ypred2 == d$y)
> ACCm2
[1] 0.542
>
> # all negative model (null model)
> ACCneg <- mean(0 == d$y)
> ACCneg
[1] 0.524
```

# Is classification accuracy sufficient?

R Code

```
> library(foreign)
> d <- read.dta("repdata.dta")
> d <- d[d$year==1982,]
> d <- na.omit(d[,c("war", "lpop", "polity2", "gdpen", "ethfrac", "lmtnest", "Oil", "cname")])
>
> d$y <- d$war
>
> crosstab(d, row.vars="y", type="frequency", percentages=F)
  y Count
  0   115
  1    21
Sum   136
> crosstab(d, row.vars="y", type="column.pct", percentages=F)
  y %
  0 0.85
  1 0.15
Sum 1.00
```

# Accuracy Paradox

R Code

```

> logit1 <- glm(war~lpop+polity2+gdpen+ethfrac+lmtnest+Oil,data=d,family= binomial(link = "logit"))
> summary(logit1)
Coefficients:
              Estimate Std. Error z value Pr(>|z|)
(Intercept) -6.38555     1.95858  -3.260  0.00111 **
lpop         0.47944     0.20212   2.372  0.01769 *
polity2      0.04639     0.04466   1.039  0.29896
gdpen       -0.32338     0.14613  -2.213  0.02690 *
ethfrac      0.77218     0.97873   0.789  0.43014
lmtnest      0.37545     0.21782   1.724  0.08476 .
Oil         -0.23107     0.77083  -0.300  0.76435
---
> d$yhat1 <- logit1$fitted
> d$ypred1 <- as.numeric(d$yhat1>.5)
>
> crosstab(d,row.vars="ypred1",col.vars="y",type="frequency",percentages=F)
      y      0      1 Sum
ypred1
0      112    19 131
1       3     2   5
Sum     115    21 136
> ACCm1 <- mean(d$ypred1 == d$y)
> ACCm1
[1] 0.8382353
>
> # all negative model (null model)
> ACCneg <- mean(0 == d$y)
> ACCneg
[1] 0.8455882

```

# Outline

- 1 Binary Outcomes
- 2 Calibration Plots
- 3 Discrimination
  - Classification accuracy
  - **Sensitivity and specificity**
- 4 Summary Accuracy Measures

# Sensitivity and specificity

	Observed	
	Positive ( $Y = 1$ )	Negative ( $Y = 0$ )
Predicted positive ( $\hat{Y} > c$ )	True Positive (TP)	False Positive (FP)
Predicted negative ( $\hat{Y} < c$ )	False Negative (FN)	True Negative (TN)
	Total Positive (P)	Total Negative (N)

- **Sensitivity:**  $TPR = TP/P$ 
  - Percent of positives that are correctly classified (also called recall or true positive rate)
  - For a screening test with higher sensitivity, fewer actual positives go undetected (lower type 2 error)
- **Specificity:**  $SPC = TN/N$ 
  - Percent of negatives that are correctly classified (also called true negative rate)
  - For a screening test with higher specificity, fewer negatives are labeled as positive (lower type 1 error)
- In theory a classifier might achieve 100% on both, but in practice there are often tradeoffs and we do not attain 100%
- The higher  $c$  the lower the sensitivity, but the higher the specificity



# Sensitivity and specificity

R Code

```
> # model 1
> crosstab(d,row.vars="ypred1",col.vars="y",type="frequency",percentages=F)
      y      0      1    Sum
ypred1
0      352    204    556
1      172    272    444
Sum     524    476   1000
> crosstab(d,row.vars="ypred1",col.vars="y",type="column.pct",percentages=F)
      y      0      1
ypred1
0      0.67 0.43
1      0.33 0.57
Sum     1.00 1.00
>
> # model 2
> crosstab(d,row.vars="ypred2",col.vars="y",type="frequency",percentages=F)
      y      0      1    Sum
ypred2
0      378    178    556
1      146    298    444
Sum     524    476   1000
> crosstab(d,row.vars="ypred2",col.vars="y",type="column.pct",percentages=F)
      y      0      1
ypred2
0      0.72 0.37
1      0.28 0.63
Sum     1.00 1.00
```

# Receiver Operating Characteristic (ROC) Curve

- ROC curves plot the sensitivity against the specificity across all cutpoints  $c$  from 0 to 1
- High discrimination ability means high sensitivity and specificity simultaneously. For such models the ROC curve goes close to the top left corner
- Low discrimination ability means low sensitivity and specificity simultaneously. For such models the ROC curve goes close to the 45 degree diagonal line.
- Can plot ROC curves to compare discrimination ability across models.


# Receiver Operating Characteristic (ROC) Curve

R Code

```
> library(plotROC)
> dd <- data.frame(y=rep(d$y,2),
+                 pred=c(d$yhat1,
+                       d$yhat2),
+                 name=rep(c("model 1","model 2"),each=nrow(d))
+                 )
>
> head(dd)
  y    pred  name
1 0 0.3030943 model 1
2 0 0.5318072 model 1
3 1 0.5119914 model 1
4 1 0.4091240 model 1
5 0 0.3163049 model 1
6 1 0.6837284 model 1

> ggroc <- ggplot(dd, aes(d = y, m = pred, color = name)) + geom_roc(labels = FALSE) + style_roc()
>
> pdf("ROC.pdf",width=9,height=6)
> ggroc
> dev.off()
```

# Receiver Operating Characteristic (ROC) Curve



# Area Under ROC Curve

- We can summarize the discrimination ability of a model by computing the **area under the estimated ROC curve (AUC)**.
- Maximum discrimination ability means  $AUC = 1$
- Minimum discrimination ability means  $AUC = 0.5$
- Can show that AUC is the probability a random pair of observations, one with  $Y = 1$  and one with  $Y = 0$ , the observation with  $Y = 1$  is correctly ranked as having higher predicted probability than the other.

# Receiver Operating Characteristic (ROC) Curve

R Code

```
> ggroc <- ggplot(dd, aes(d = y, m = pred, color = name)) + geom_roc(labels = FALSE) + style_roc()
>
> calc_auc(ggroc)
  PANEL group      AUC
1      1     1 0.6654632
2      1     2 0.7413461

> library(pROC)
> roccurve <- roc(d$y ~ d$yhat1)
> auc(roccurve)
Area under the curve: 0.6655
```

# Outline

- 1 Binary Outcomes
- 2 Calibration Plots
- 3 Discrimination
  - Classification accuracy
  - Sensitivity and specificity
- 4 Summary Accuracy Measures

# Brier Score

To summarize the accuracy of binary classifiers we use the Brier score:

$$BS = \frac{1}{N} \sum_i^N (\hat{Y}_i - Y_i)^2$$

- If  $\hat{Y}_i = 1$  and  $Y_i = 1$ , then  $BS = 0$ , the best score achievable.
- If  $\hat{Y}_i = 1$  and  $Y_i = 0$ , then  $BS = 1$ , the worst score achievable.
- Note that this measure is based on relative deviations ( $\hat{Y}_i = .51$  pays a much larger cost than  $\hat{Y}_i = .99$  for  $Y_i = 1$ ), rather than binary deviations (such as correctly or incorrectly classified for a given cutoff  $c$ )
- Note that this is proper score function:
  - If  $\hat{Y}$  is the vector of predictions and you use a scoring function to give a reward of  $S(\hat{Y}, i)$  if the  $i$ -th event occurs
  - The highest expected reward is obtained by reporting the true probability distribution so you encourages the forecaster to be honest to maximize the expected reward



# Brier Score

BS can be decomposed as the sum of a measure of calibration and refinement:

$$BS = \frac{1}{N} \sum_i^N (\hat{Y}_i - Y_i)^2 = \frac{1}{N} \sum_k^K (\hat{Y}_k - \bar{Y}_k)^2 + \frac{1}{N} \sum_k^K n_k (\bar{Y}_k (1 - \bar{Y}_k))$$

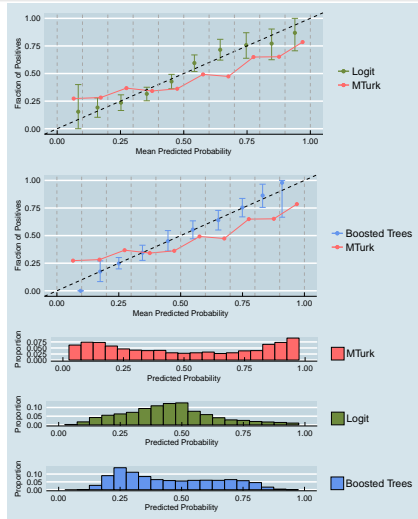
- $K$  is the # of unique predicted probabilities,  $n_k$  is the # of predictions in predicted probability category  $k$  (s.t.  $N = \sum_k n_k$ ),  $\bar{Y}_k$  is the observed frequency in  $k$  given the predicted probability  $\hat{Y}_k$
- The first term captures calibration
  - How much do observed frequency of positives in  $k$  vary from predicted fraction  $\hat{Y}_k$
- The second term refinement (related to the area under the ROC curve)
  - How uniform is occurrence of outcomes within each group  $k$

# Brier Score

R Code

```
> mean((d$ypred1-d$y)^2)
[1] 0.1617647
> mean((0-d$y)^2)
[1] 0.1544118
```

# Predicting Recidivism



Brier scores are: Mturk .24, Logit = .21 and Boosted Trees = .19 (all three model are similar on AUC)

# Cross-Validation

- Often it is useful to validate model looking at out of sample fit instead of in-sample fit
- Most common method is  $k$ -fold Cross Validation
- Cut the dataset randomly into say 10 folds
- Start a loop, pick the first fold as the hold out, fit the model on the remaining 9 folds and predict outcomes for hold out fold
- Repeat while using next fold as hold out
- In the end you have an out of sample prediction for each observation and then examine model accuracy using this prediction
- Common is 10 folds, but can use other values

# Cross-Validation

R Code

```

> ## bail example cross validation
> rm(list=ls())
> d <- read.csv("bail.csv")
> d$y <- d$two_year_recid
>
> Foldsize <- nrow(d)/10
> Foldsize
[1] 100
>
> # fold var
> d$fold <- sample(rep(1:10,each=Foldsize),nrow(d),replace=F)
>
> d$yhat1 <- NA
> d$yhat2 <- NA
>
> for(i in 1:10){
+
+ # model 1
+ logit1 <- glm(y~sex+age,data=d[d$fold!=i,])
+ d$yhat1[d$fold==i] <- predict(logit1,newdata=d[d$fold==i,])
+
+ # model 2
+ logit2 <- glm(y~sex+age+log(1+priors_count),data=d[d$fold!=i,])
+ d$yhat2[d$fold==i] <- predict(logit2,newdata=d[d$fold==i,])
+
+ }

```

# Cross-Validation


R Code

```

> dd <- data.frame(y=rep(d$y,2),
+                 pred=c(d$yhat1,
+                       d$yhat2),
+                 name=rep(c("model 1", "model 2"),each=nrow(d))
+ )
>
> head(dd)
  y      pred  name
1 0 0.2103555 model 1
2 0 0.4504963 model 1
> dd$name <- factor(dd$name)
>
> ggroc <- ggplot(dd, aes(d = y, m = pred, color = name)) + geom_roc(labels = FALSE) + style_roc()
> pdf("ROC2.pdf",width=9,height=6)
> ggroc
> dev.off()
null device
      1
>
> # AUC
> calc_auc(ggroc)
  PANEL group      AUC
1     1     1 0.6183747
2     1     2 0.7311746
>
> # BS
> mean((d$yhat1-d$y)^2)
[1] 0.2383642
> mean((d$yhat2-d$y)^2)
[1] 0.2106819

```

# Cross-Validation ROC Curve



# Cross-Validation

R Code

```

> logit2      <- glm(y~sex+age+log(1+priors_count),data=d)
> d$yhat2in   <- predict(logit2,newdata=d)
>
> dd <- data.frame(y=rep(d$y,2),
+                 pred=c(d$yhat2in,
+                       d$yhat2),
+                 name=rep(c("model 1","model 2"),each=nrow(d))
+ )
> ggroc <- ggplot(dd, aes(d = y, m = pred, color = name)) + geom_roc(labels = FALSE) + style_roc()

> # AUC
> calc_auc(ggroc)
  PANEL group      AUC
1      1      1 0.7369259
2      1      2 0.7311746
>
>
> # BS
> mean((d$yhat2in-d$y)^2)
[1] 0.2087114
> mean((d$yhat2-d$y)^2)
[1] 0.2106819

```



# Cross-Validation ROC Curve

